

Base and Bounds (contd)

Involvement of OS

The OS must take action to implement base and bound approach

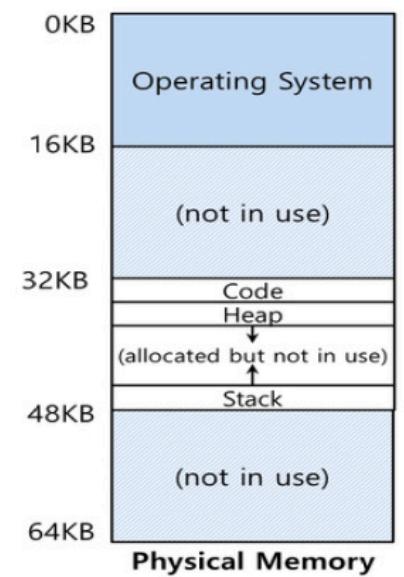
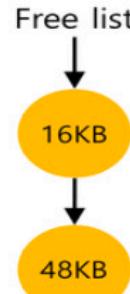
- Three critical junctures →
- * when a process is created
 - * when a process is terminated
 - * when context switch occurs

When a process is created →
OS must be able to allocate
memory to process,

keep track which parts of free
memory are not in use.

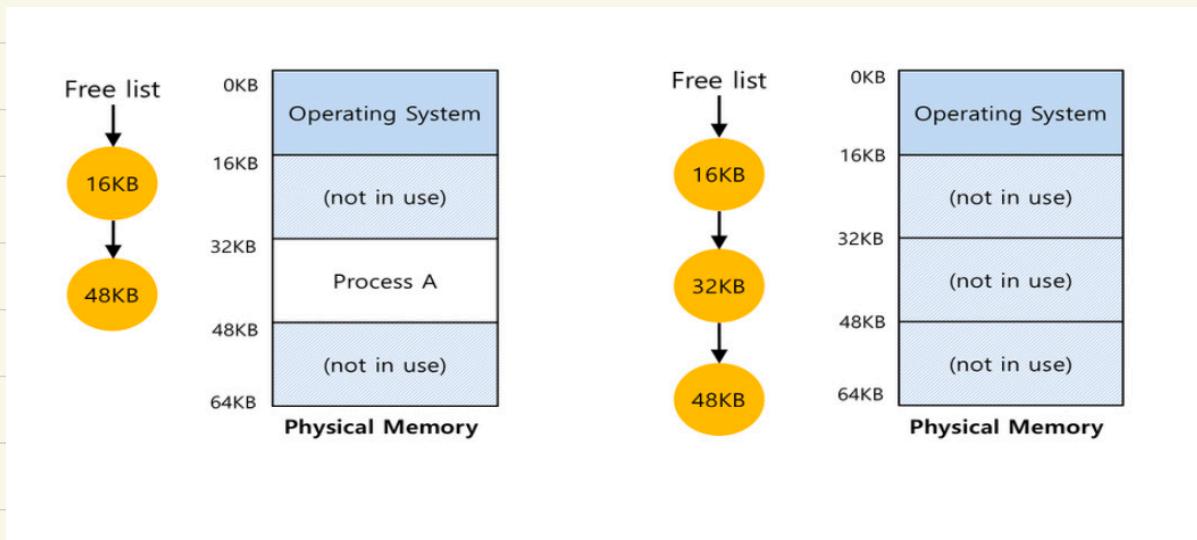
A free list is a list of the ranges
of the physical memory which
are not currently in use.

The OS looks up the free list

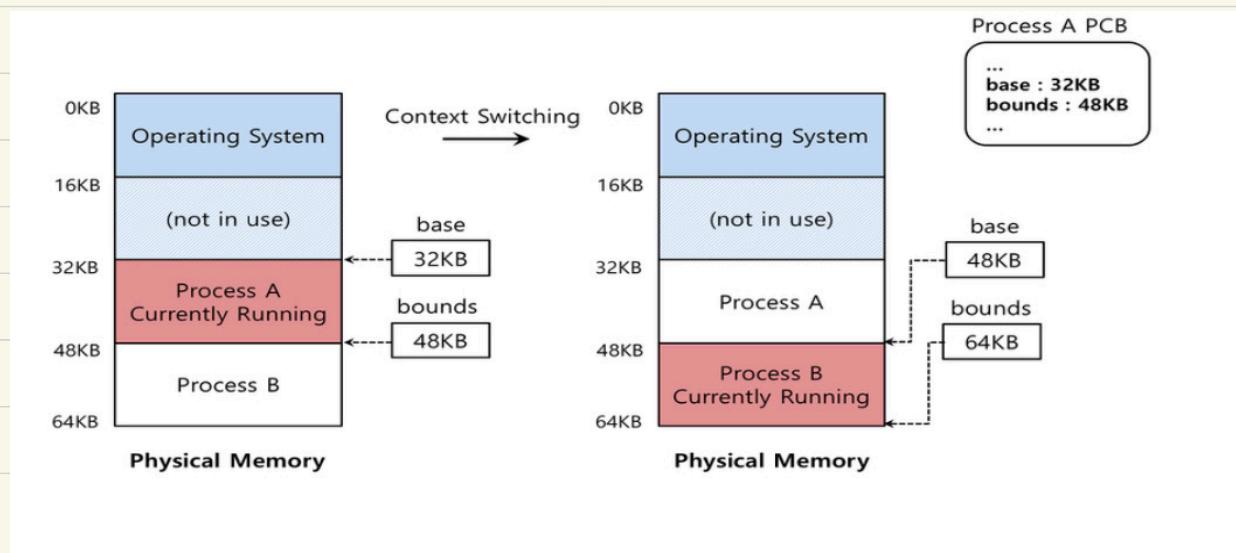


when a process is terminated

→ the OS must put the memory back on the free list.



when context switch occurs → the OS must save and restore the base and bounds pair in the process control block (PCB)



Apart from above, OS must provide exception handlers.

- * if a process tries to access memory outside its bounds · the CPU will raise an exception
- * the OS must be prepared to take action when such exception arises.

Hardware/OS interaction in timeline

OS @ run (kernel mode)	Hardware	Program (user mode)
To start process A: allocate entry in process table alloc memory for process set base/bound registers return-from-trap (into A)	restore registers of A move to user mode jump to A's (initial) PC	Process A runs Fetch instruction
	translate virtual address perform fetch	Execute instruction
	if explicit load/store: ensure address is legal translate virtual address perform load/store	
	Timer interrupt move to kernel mode jump to handler	(A runs...)
Handle timer decide: stop A, run B call switch() routine save regs(A) to proc-struct(A) (including base/bounds) restore regs(B) from proc-struct(B) (including base/bounds) return-from-trap (into B)	restore registers of B move to user mode jump to B's PC	Process B runs Execute bad load
	Load is out-of-bounds; move to kernel mode jump to trap handler	
Handle the trap decide to kill process B deallocate B's memory free B's entry in process table		

Pros of Base and Bound approach

- * simple
- * fast (2 registers, adder, comparator)

Cons of Base and Bound approach

- * can't keep programs from accidentally overwriting its own code

(because, here CPU only ensures you stay within your own allocated region)

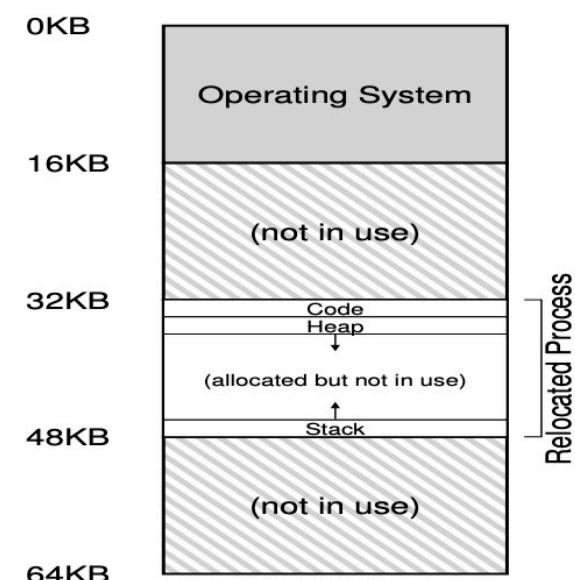
- * can't share code/data with other processes.

(because, the mapping is contiguous and exclusive)

- * can't grow stack/heap as needed.

- * Internal fragmentation

if the process's stack and heap are not too big,
all the space b/w the two is simply wasted.



Question: How do we support a large address space with (potentially) a lot of free space b/w the stack and the heap?

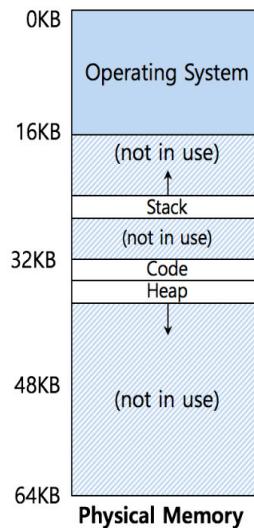
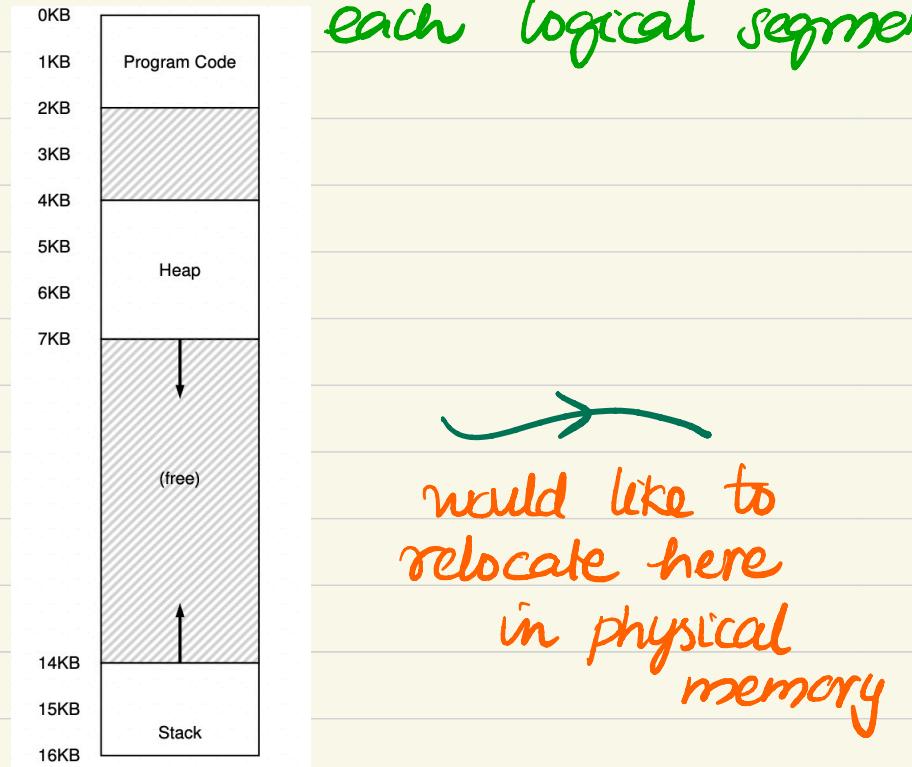
Segmentation

Segment is just a contiguous portion of the address space of a particular length.

- logically different segments: code, heap, stack

each segment can be placed in different part of physical memory.

How? Instead of one base and bounds pair, maintain a pair for each logical segments, code, heap and stack.

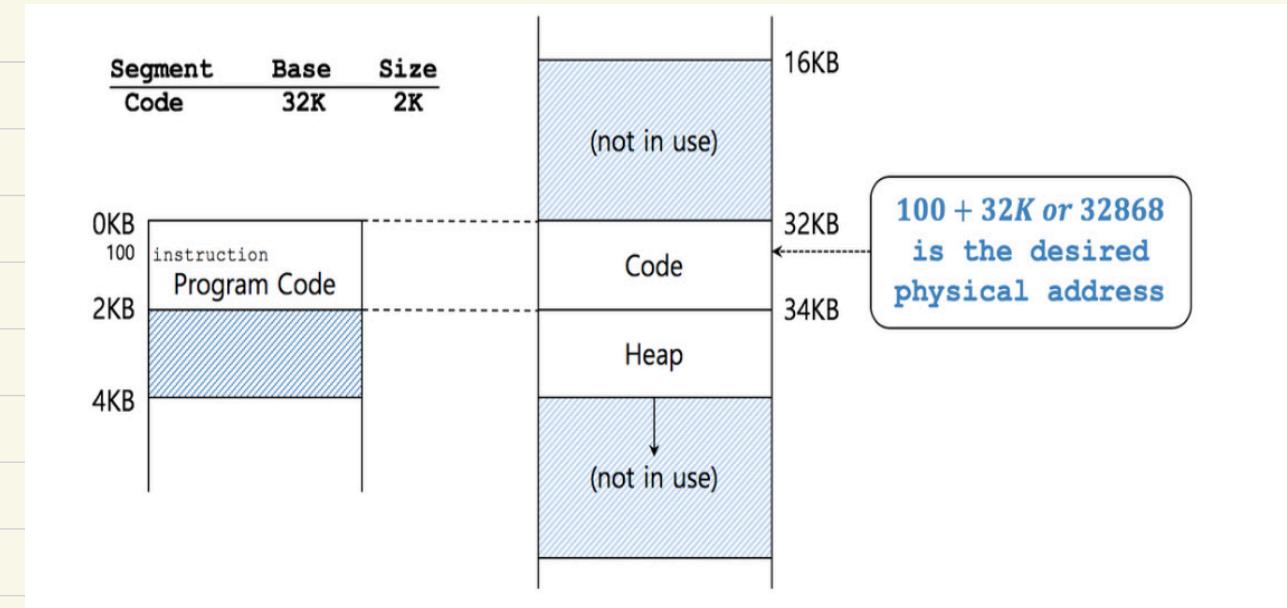


Segment	Base	Size
Code	32K	2K
Heap	34K	2K
Stack	28K	2K

only used memory in each segment is allocated a space, as specified by this size.

Virtual address space

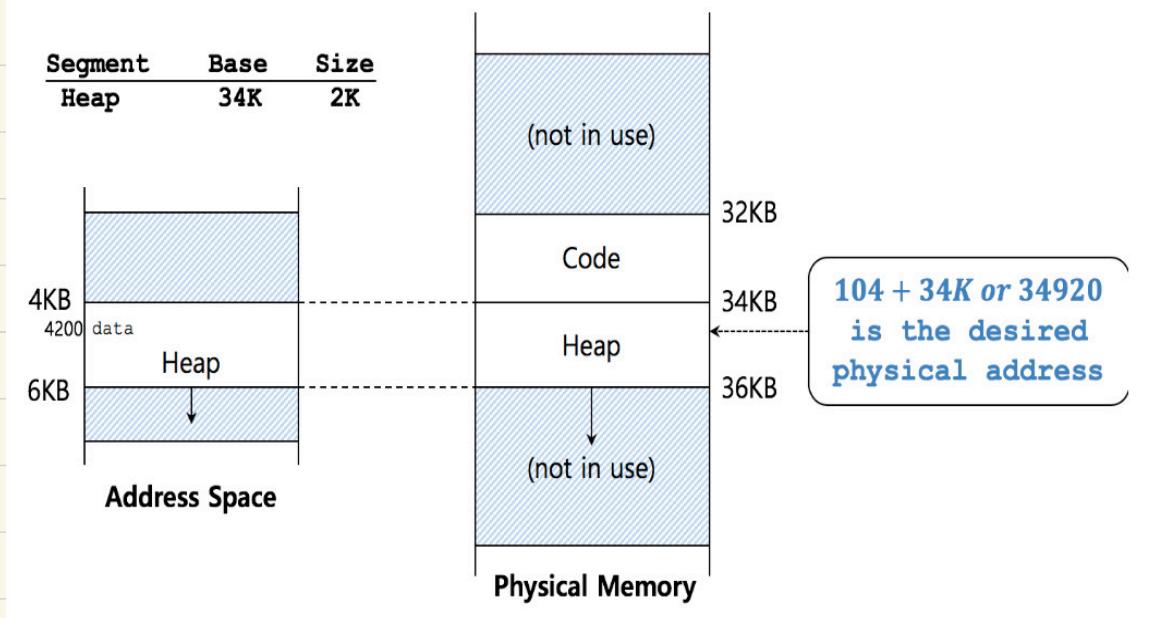
eg:- Suppose that a reference is made to virtual address 100
 → virtual address - start address of the ^(code segment)
 * the offset of virtual address 100 is 100. (the code segment starts at virtual address zero)
 $\Rightarrow \text{physical address} = \text{offset} + \text{base}$



* Now, lets look at an address in the heap, virtual address 4200
 (recall, base of heap = 34K)
 $\text{virtual address} + \text{Base} = 4200 + 34816 = 39016 \times \text{not correct!}$

what to do? Extract the offset into the heap.

Note that heap starts at 4KB \Rightarrow offset = 4200 - (4 × 1024) = 104

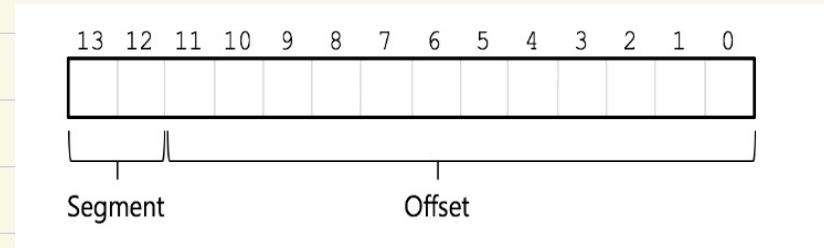


Remark: If an illegal address such as 7KB, which is beyond the end of heap is referenced, the hardware detects that the address is out of bounds \rightarrow traps into the OS \rightarrow leads to the famous "segmentation fault"

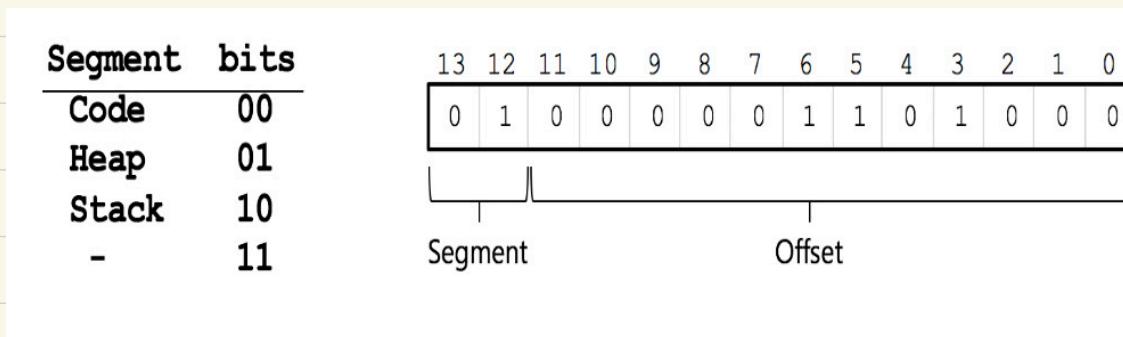
Question: While translating, how does the hardware know the offset into a segment, and to which segment an address refers?

* Explicit approach

Chop up the address space into segments based on the top few bits of virtual address.



e.g.: - virtual address 4200 (01000001101000)



→ 0000 1101 1000
= 104 in decimal

* Bound check is simple with offset
(just need to check, whether offset ≤ bounds)

```

1 // get top 2 bits of 14-bit VA
2 Segment = (VirtualAddress & SEG_MASK) >> SEG_SHIFT
3 // now get offset
4 Offset = VirtualAddress & OFFSET_MASK
5 if (Offset >= Bounds[Segment])
6     RaiseException(PROTECTION_FAULT)
7 else
8     PhysAddr = Base[Segment] + Offset
9     Register = AccessMemory(PhysAddr)

```

* SEG-MASK = 110000 0000 0000

* SEG-SHIFT = 12 * OFFSET-MASK = 00 1111 1111 1111

Disadvantages of this approach

* top two bits can be actually used for four segments, and we have only three segments \Rightarrow virtual address space is under utilized.

(a way to avoid this problem; by putting code and heap in a single segment and stack in the other, so that only 1 bit is been used for 2 segments)

* each segment is limited to maximum size (eg:- in the above case 4KB)

* Implicit approach

The hardware determines the segment by noticing how the address was formed.

* If the address was generated from the program counter (\Rightarrow it was an instruction fetch)
 \Rightarrow the address is within the code segment.

* If the address is based off of the stack or base pointer,
it must be in the stack segment.

* any other address must be in the heap.

Question: How to refer the stack segment?

What is special about stack segment? \rightarrow it grows backward.

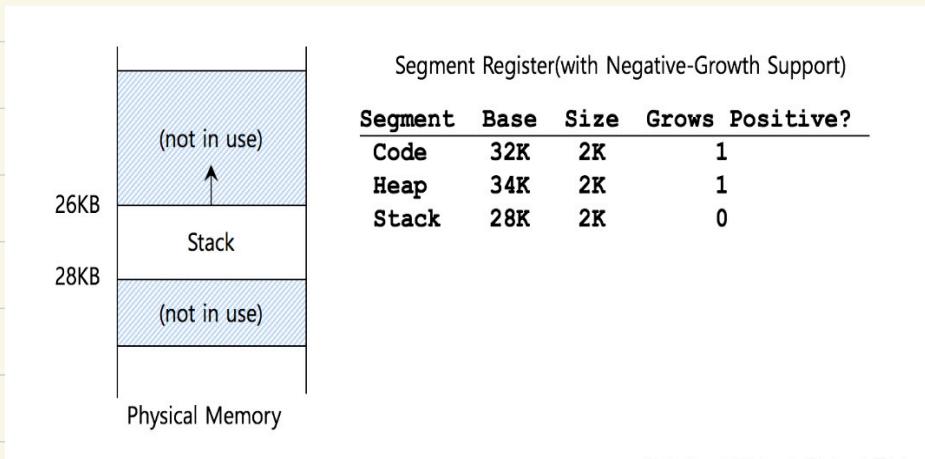
For instance, in the previous example, in physical memory, stack starts at 28KB and grows back to 26KB (corr. to VA, 16KB to 14KB)

\therefore we need a mechanism to know the direction of the growth

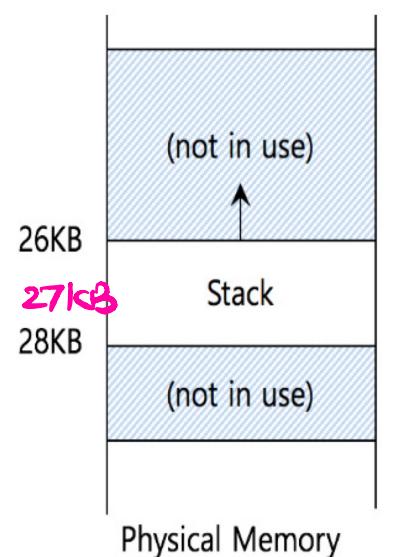
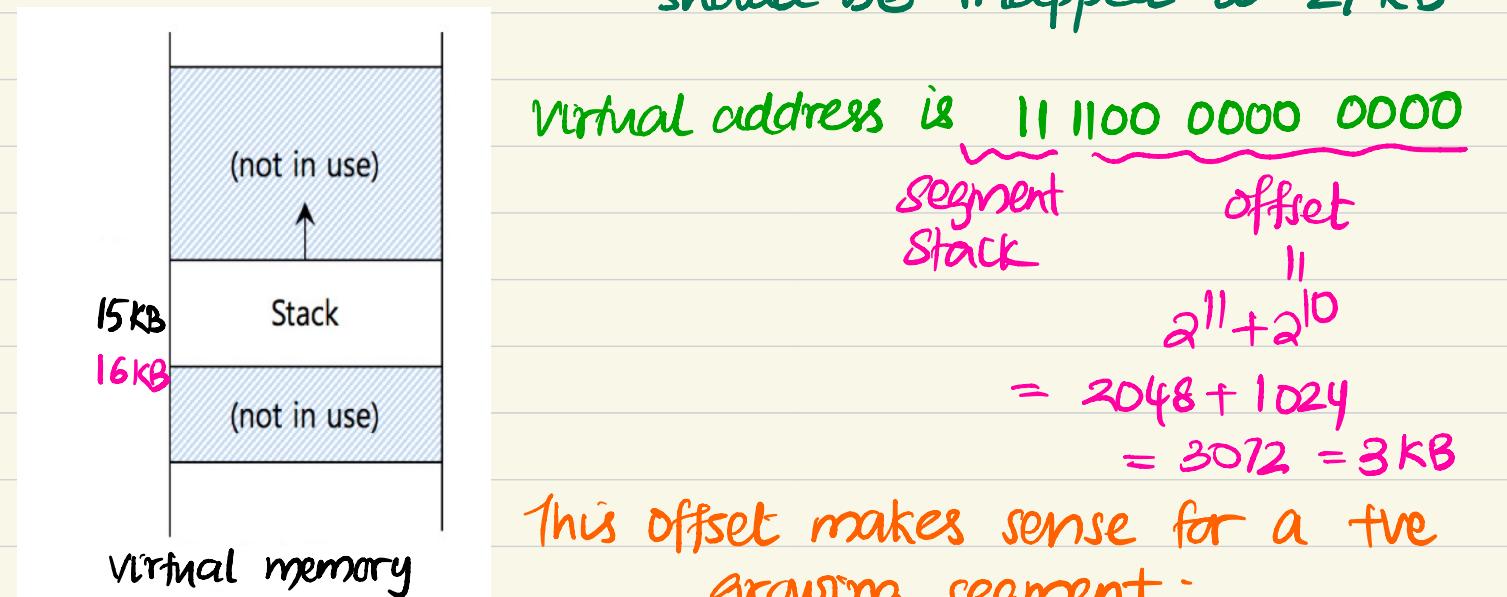
\hookrightarrow extra hardware support!

* the hardware checks which way the segment grows.

* 1: positive direction 0: negative direction.



consider an example, assume we wish to access virtual address 15KB, which should be mapped to 27 KB.



But here, the segment grows downwards from a high address toward lower addresses. In hardware's view, the base (28 kB) is the top (highest address) of the segment, and the offset should decrease as you go deeper in it.

⇒ we cannot add the above offset 3kB to the virtual address!

∴ To obtain the correct -ve offset (-1 kB), we must subtract the maximum segment size from 3kB.

Here, max segment size = 4 kB

⇒ correct -ve offset = +ve offset - max_seg = -1 kB (i.e., 1 kB below the base address in memory)

$$\begin{aligned}\Rightarrow \text{physical address} &= \text{Base} + (\text{-ve}) \text{ offset} \\ &= 28\text{ kB} + -1\text{ kB} = 27\text{ kB}\end{aligned}$$

Modified algorithm

Segment = (virtual address & SEG-MASK) \gg SEG-SHIFT

Offset = virtual address & OFFSET-MASK

If (Direction[segment] == 1) {

if (offset ≥ Bounds[segment])

Raise Exception (PROTECTION-FAULT)

else

PhysAddr = Base[segment] + offset }

(Because of the 2kB bound, this should go only ≤ 2kB below)

else {

Disp = offset - MaxSize[segment]

if (EDISP) ≥ Bounds[segment])

Raise Exception (PROTECTION FAULT)

else

PhysAddr = Base[segment] + Disp

}