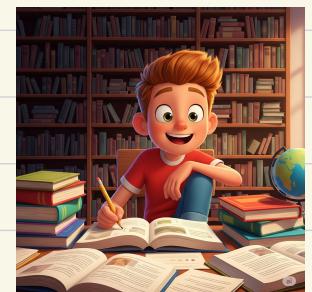


Paging: Faster translations (TLBs)

Recall the problem,

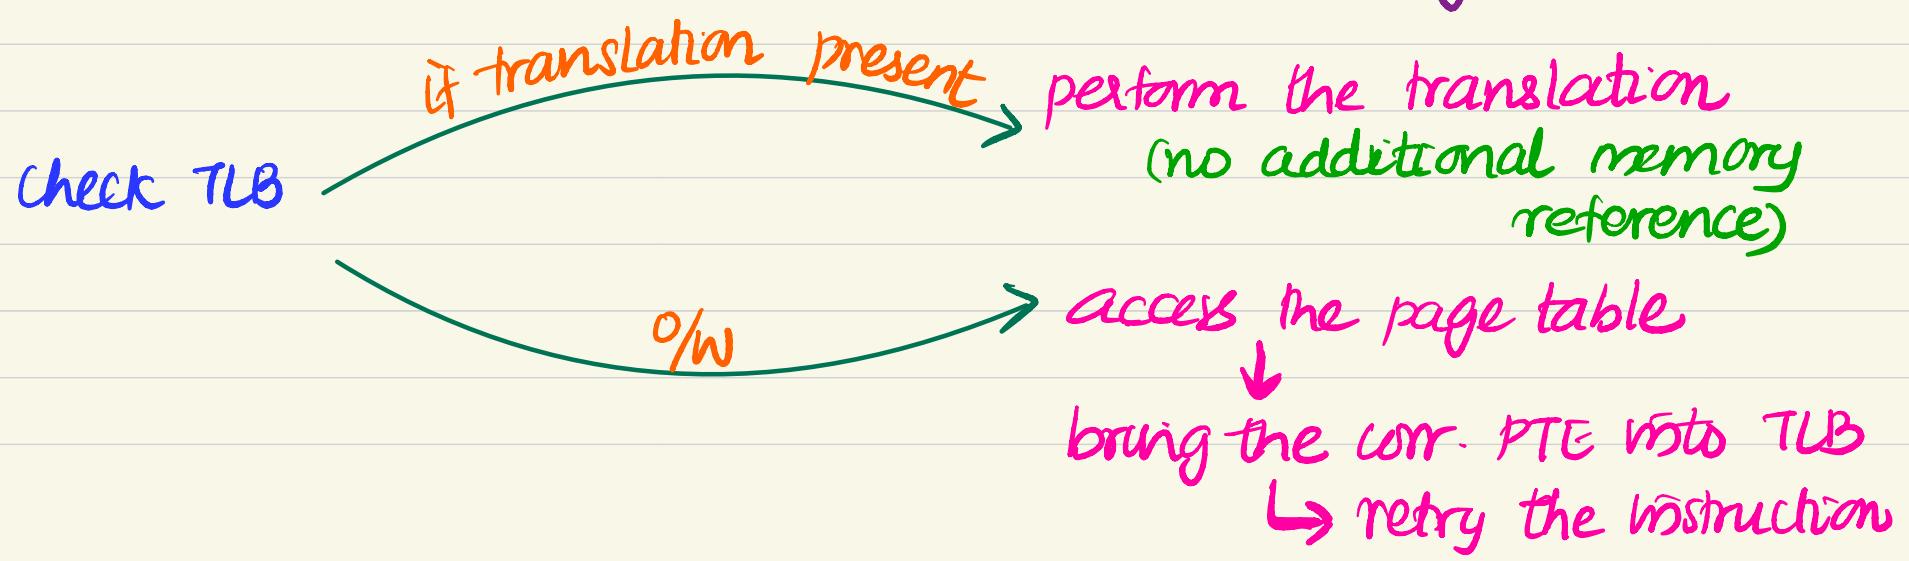
Question: How to speed up address translation and generally avoid the extra memory reference?



↳ using hardware support
"Translation lookaside Buffer (TLB)"
today's focus!

A TLB is a part of chip's Memory-management unit (MMU), and is simply a cache of popular virtual-to-physical address translations.

For each memory reference, hardware do the following.



```

1  VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2  (Success, TlbEntry) = TLB_Lookup(VPN)
3  if (Success == True) // TLB Hit
4      if (CanAccess(TlbEntry.ProtectBits) == True)
5          Offset = VirtualAddress & OFFSET_MASK
6          PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7          Register = AccessMemory(PhysAddr)
8      else
9          RaiseException(PROTECTION_FAULT)
10     else // TLB Miss
11         PTEAddr = PTBR + (VPN * sizeof(PTE))
12         PTE = AccessMemory(PTEAddr)
13         if (PTE.Valid == False)
14             RaiseException(SEGMENTATION_FAULT)
15         else if (CanAccess(PTE.ProtectBits) == False)
16             RaiseException(PROTECTION_FAULT)
17         else
18             TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
19             RetryInstruction()

```

TLB control flow algorithm of a hardware-managed TLB .

Idea behind TLB: if the address found in TLB → TLB hit -
 (less overhead)
 O/W → TLB miss
 (extra memory reference needed)

Goal: avoid TLB misses as much as possible.

example: Assume that we have an array of 10 4-byte integers in memory starting at virtual address 100.

Assumptions: virtual address space \rightarrow 8 bits \rightarrow 256 bytes.
page size \rightarrow 16 bytes \Rightarrow offset = 4 bits

$$\Rightarrow \# \text{ pages} = \frac{2^8}{2^4} = 2^4 = 16 \text{ pages}$$

bits for VPN = 4 bits

	Offset				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 02					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06	a[0]	a[1]	a[2]		
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

$a[0] \rightarrow VA = 100 \rightarrow \underline{0110}, \underline{0100}$

VPN = 6 offset = 4

Consider the foll. code that access each array element.

```
int i, sum=0;  
for (i=0 ; i<10; i++) {  
    sum+=a[i]  
}
```

first access: $a[0] \rightarrow VA 100$ (VPN = 6, offset = 4)

Check TLB \rightarrow TLB miss \rightarrow Insert in TLB (Translation of VPN 6)

second access \rightarrow $a[1]$: VPN = 6, offset = 8

Check TLB \rightarrow TLB hit

||^{by} $a[2] \rightarrow$ hit $a[3] \rightarrow$ miss $a[4], a[5], a[6] \rightarrow$ hit
(Translation of VPN 7)
 $a[7] \rightarrow$ miss $a[8], a[9] \rightarrow$ hit.
(Translation of VPN 8)

Summary: out of 10 accesses \rightarrow 7 hits and 3 misses

\Rightarrow 70% TLB hit rate

Why high TLB hit rate?

* spatial locality: A program accessing elements of the memory which are placed close to each other.

\Rightarrow only the first access to an element in a page yields a TLB miss.

\Rightarrow larger size \rightarrow high hit rate.

a typical page size is like, 4KB. In that case, even if the array size

is 2048, we might end up in having only 2 TLB misses (if a[0] is at the starting of the first page) \Rightarrow Hit rate = $\frac{2046}{2048} \times 100 = 99.9\%$, which is pretty good!

* Temporal locality: if the program accesses the same memory locations frequently, then it will result in fewer TLB misses.

So, if the program has the above two properties then we have high TLB hit rates.

What happens in case of TLB miss?

who handles a TLB miss

→ Hardware managed TLB
eg:- Intel x86 architecture

on a TLB miss, hardware only raises an exception ↗

→ Software managed TLB
eg:- MIPS/SPARC/RISC-V

Exception happens \rightarrow current execution is paused \rightarrow kernel mode



retry instruction \leftarrow return-from-trap \leftarrow update TLB \leftarrow trap handler

here, slightly different from usual return-from-trap as the execution needs to resume from the instruction that caused the trap (not the next one)

```

1  VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2  (Success, TlbEntry) = TLB_Lookup(VPN)
3  if (Success == True) // TLB Hit
4      if (CanAccess(TlbEntry.ProtectBits) == True)
5          Offset = VirtualAddress & OFFSET_MASK
6          PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7          Register = AccessMemory(PhysAddr)
8      else
9          RaiseException(PROTECTION_FAULT)
10 else // TLB Miss
11     RaiseException(TLB_MISS)

```

TLB control flow algorithm of an OS handled TLB

Hardware handled TLB

Software handled TLB

How it works?

On a TLB miss, the CPU hardware (page table walker) walks the page tables in memory, finds the right PTE and inserts into TLB.

(needed in multi-level page tables)
which we will see later)

- fast miss handling

(no OS trap/interrupt needed)

on a TLB miss, the CPU traps into the OS, and the OS trap handler explicitly finds the PTE in memory and writes the entry into the TLB.

Advantages

- simplicity of hardware

(CPU doesn't need a page table walker)

- transparent to OS
(OS just maintains page tables, hardware does the rest)
- flexibility for OS
(OS can define page-table format, replacement policy, etc.)

Disadvantages

- Hardware complexity
(needs dedicated logic (like, page-walker) increasing CPU design cost and power)
- less flexibility
(OS can't easily change TLB replacement policy or page-table format - it must follow the hardware defined scheme)
- high miss penalty
TLB miss → trap → OS handler
→ insert entry → return to user
(slower than hardware)
- OS burden

Note that, in software handled TLB systems, (like, MIPS), the OS has full freedom to design its own mapping structures.

- linear page tables, inverted page tables, hash-based structures, etc..

no fixed page table format ⇒ no point in a base register!

Caution! The trap handler should not cause an infinite chain of TLB misses!

How this happens? Imagine a system where:

virtual address → needs translation → TLB miss → go to page table.

But page table itself is stored at some virtual address (not physical)
- then to access PTE, the OS/hardware tries to translate the page table's
virtual address. That lookup itself causes a table miss, and
the chain continues forever---

How OS avoids the possibility of infinite recursion?

- it may keep TLB miss handlers in physical memory
(where they are unmapped and not subject to address translation)
- or reserve some entries in the TLB for permanently-valid
translations and use some of those permanent translation slots for
the handler code itself (these wired translations always hit in the TLB).